

Modelling Cryptocurrency exchange rates using ARMA Models

Applied Macroeconometrics Assignment I

Rishabh Patil | 2021A7PS0464H

Contents

Description of Dataset	3
The data	3
Candlestick Chart (daily pricing)	3
Line Graph(daily pricing)	4
Clustering Month-wise avg price:	5
Line Graph	6
Candlestick Graph	6
Returns	7
Visual Interpretation	9
Estimation and Holdback dataset	11
Model fitting and model selection	15
Analyzing ACF and PACF	15
Testing ARMA models	17
AR(1)	17
Coefficients Test	18
Residual Check	18
Stationarity, Invertibility and Causality	19
Mean of the model	20
ARMA(2,1)	20
Coefficients Test	21
Residual Check	22
Stationarity, Invertibility and Causality	23
Mean	23
ARMA(1,2)	23
Coefficients Test	24
Residual Check	24

Stationarity, Invertibility and Causality	25
Mean	25
MA(2)	25
Coefficients Test	26
Residual Check	27
Stationarity, Causality and Invertibility	28
Mean	28
Ljung-Box test for AR(1) MA(2) and ARMA(1,2)	28
AIC	32
Forecasting	32
Paired F-test	42
AR1 ARMA192	42
AR1 ARMA(2,19)	43
ARMA(19,2) ARMA(2,19)	43
DM Test	43
Absolute loss function	44
Table 2	45

```
knitr::opts_chunk$set(tidy.opts = list(width.cutoff = 60), tidy = TRUE)
```

```
# install.packages('quantmod')
```

```
library(quantmod)
```

```
## Loading required package: xts
```

```
## Loading required package: zoo
```

```
##
```

```
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## as.Date, as.Date.numeric
```

```
## Loading required package: TTR
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
## method from
```

```
## as.zoo.data.frame zoo
```

Description of Dataset

For our study, we'd be analyzing the exchange rates of the cryptocurrency Bitcoin, more precisely the weekly returns.

The variable of concern here is the price(in USD) of Bitcoin (returns on it).

Since there is a general notion of cryptocurrencies having a 4 year cycle, we'd be looking at 8-10 year data. And to make our analysis feasible, instead of a daily chart, we would be analyzing a weekly rolling average for our time period of 10 years.

The data

```
btc_df <- getSymbols("BTC-USD", src = "yahoo", from = Sys.Date() -  
  365 * 10, to = Sys.Date(), auto.assign = FALSE)  
print(head(btc_df))
```

```
##          BTC-USD.Open BTC-USD.High BTC-USD.Low BTC-USD.Close BTC-USD.Volume  
## 2014-09-17      465.864      468.174      452.422      457.334      21056800  
## 2014-09-18      456.860      456.860      413.104      424.440      34483200  
## 2014-09-19      424.103      427.835      384.532      394.796      37919700  
## 2014-09-20      394.673      423.296      389.883      408.904      36863600  
## 2014-09-21      408.085      412.426      393.181      398.821      26580100  
## 2014-09-22      399.100      406.916      397.130      402.152      24127600  
##          BTC-USD.Adjusted  
## 2014-09-17          457.334  
## 2014-09-18          424.440  
## 2014-09-19          394.796  
## 2014-09-20          408.904  
## 2014-09-21          398.821  
## 2014-09-22          402.152
```

Candlestick Chart (daily pricing)

```
chartSeries(btc_df, name = "BTC-USD", subset = "last 120 months")
```

```
## Warning in last.xts(structure(c(465.864013671875, 456.859985351562,  
## 424.102996826172, : requested length is greater than original
```




Line Graph(daily pricing)

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 4.2.3
```

```
ggplot(data = btc_df, aes(y = btc_df$`BTC-USD.Close`, x = index(btc_df)),
  group = 1) + geom_line()
```

```
## Don't know how to automatically pick scale for object of type <xts/zoo>.
## Defaulting to continuous.
```



Clustering Month-wise avg price:

```
# install.packages('xts')
```


```
library(xts)
btc_week_df <- apply.weekly(btc_df, FUN = mean)
head(btc_week_df)
```

```
##          BTC-USD.Open BTC-USD.High BTC-USD.Low BTC-USD.Close BTC-USD.Volume
## 2014-09-21      429.9170      437.7182      406.6244      416.8590      31380680
## 2014-09-28      410.6507      418.6690      399.3771      407.6926      26681800
## 2014-10-05      369.7743      376.7210      353.2071      361.4266      39522557
## 2014-10-12      346.9274      363.3089      337.5679      355.2346      48736115
## 2014-10-19      389.0103      397.7904      380.4106      390.4799      22414581
## 2014-10-26      372.2030      377.1116      362.5564      367.3164      16241686
##          BTC-USD.Adjusted
## 2014-09-21      416.8590
## 2014-09-28      407.6926
## 2014-10-05      361.4266
## 2014-10-12      355.2346
## 2014-10-19      390.4799
## 2014-10-26      367.3164
```

Line Graph


```
ggplot(data = btc_week_df, aes(y = btc_week_df$`BTC-USD.Close`,  
  x = index(btc_week_df)), group = 1) + geom_line()
```

```
## Don't know how to automatically pick scale for object of type <xts/zoo>.  
## Defaulting to continuous.
```



Candlestick Graph

```
chartSeries(btc_week_df)
```



Returns


For plotting the returns, we use this formula:

$$R = \frac{S(t+1) - S(t)}{S(t)} \text{ or plotting the log return: } RLt = \ln\left(\frac{S(t+1)}{S(t)}\right)$$

```
daily_returns <- log(btc_df/lag(btc_df, 1))
ggplot(data = daily_returns, aes(y = daily_returns$`BTC-USD.Close`,
  x = index(daily_returns)), group = 1) + geom_line()
```

```
## Don't know how to automatically pick scale for object of type <xts/zoo>.
## Defaulting to continuous.
```


```
## Warning: Removed 1 row containing missing values (`geom_line()`).
```



```
weekly_return <- log(btc_week_df$`BTC-USD.Close`/lag(btc_week_df$`BTC-USD.Close`,
1))
ggplot(data = weekly_return, aes(y = weekly_return$`BTC-USD.Close`,
x = index(weekly_return)), group = 1) + geom_line()
```

```
## Don't know how to automatically pick scale for object of type <xts/zoo>.
## Defaulting to continuous.
```

```
## Warning: Removed 1 row containing missing values (`geom_line()`).
```

This is the basis of our analysis

Visual Interpretation

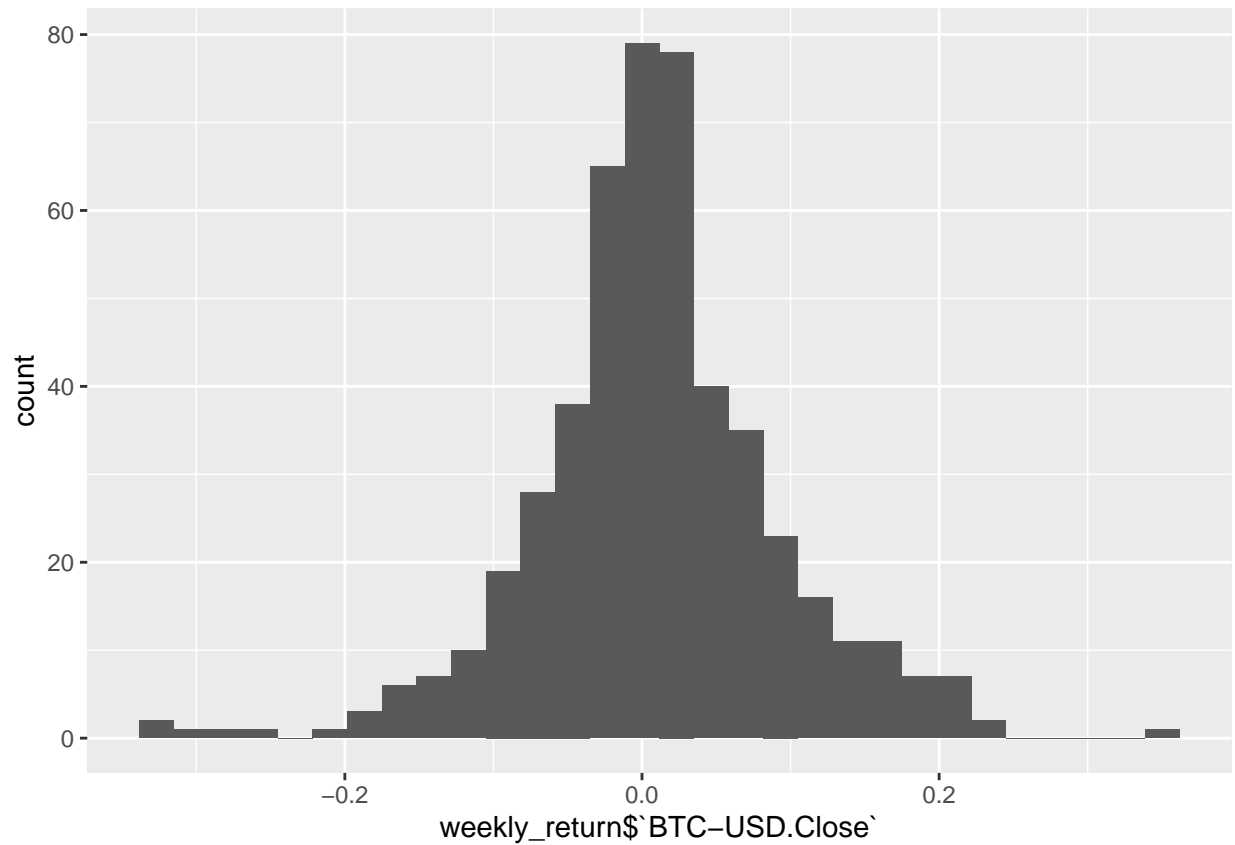
From the above log returns chart we can see that the long run expectation for weekly return is 0

$$\mathbb{E}[y_t] = 0$$

```
ggplot(weekly_return, aes(weekly_return$`BTC-USD.Close`)) + geom_histogram()
```

```
## Don't know how to automatically pick scale for object of type <xts/zoo>.
## Defaulting to continuous.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

## Warning: Removed 1 rows containing non-finite values (`stat_bin()`).
```




The data looks normally distributed with slight outliers.

looking for seasonality:

```
yearly_return <- yearlyReturn(btc_df)
ggplot(data = yearly_return, aes(y = yearly_return$yearly.returns,
  x = index(yearly_return), group = 1) + geom_line()
```

```
## Don't know how to automatically pick scale for object of type <xts/zoo>.
## Defaulting to continuous.
```



We do see that there is a 4-year cycle apparent from the yearly returns, that is, the returns rise up for a year, then reach peak very steeply the next and a steep decline sets in ending the cycle.

But apart from the yearly return analysis the weekly returns appear stationary.

The objective of this study is to model the weekly returns with appropriate ARMA model(s) and predict the trend (forecast it).

Estimation and Holdback dataset

Estimation: from 2014-09-21 to 2022-09-21 - *train data*


Holdback: from 2022-09-28 to 2024-02-23 - *test data*

(~75/25 split) (split post cycle)

```
btc_week_est <- btc_week_df[index(btc_week_df) >= as.Date("2014-09-21") &
  index(btc_week_df) <= as.Date("2022-09-21"), ]

btc_week_hb <- btc_week_df[index(btc_week_df) >= as.Date("2022-09-22") &
  index(btc_week_df) <= as.Date("2024-02-23"), ]

chartSeries(btc_week_est)
```



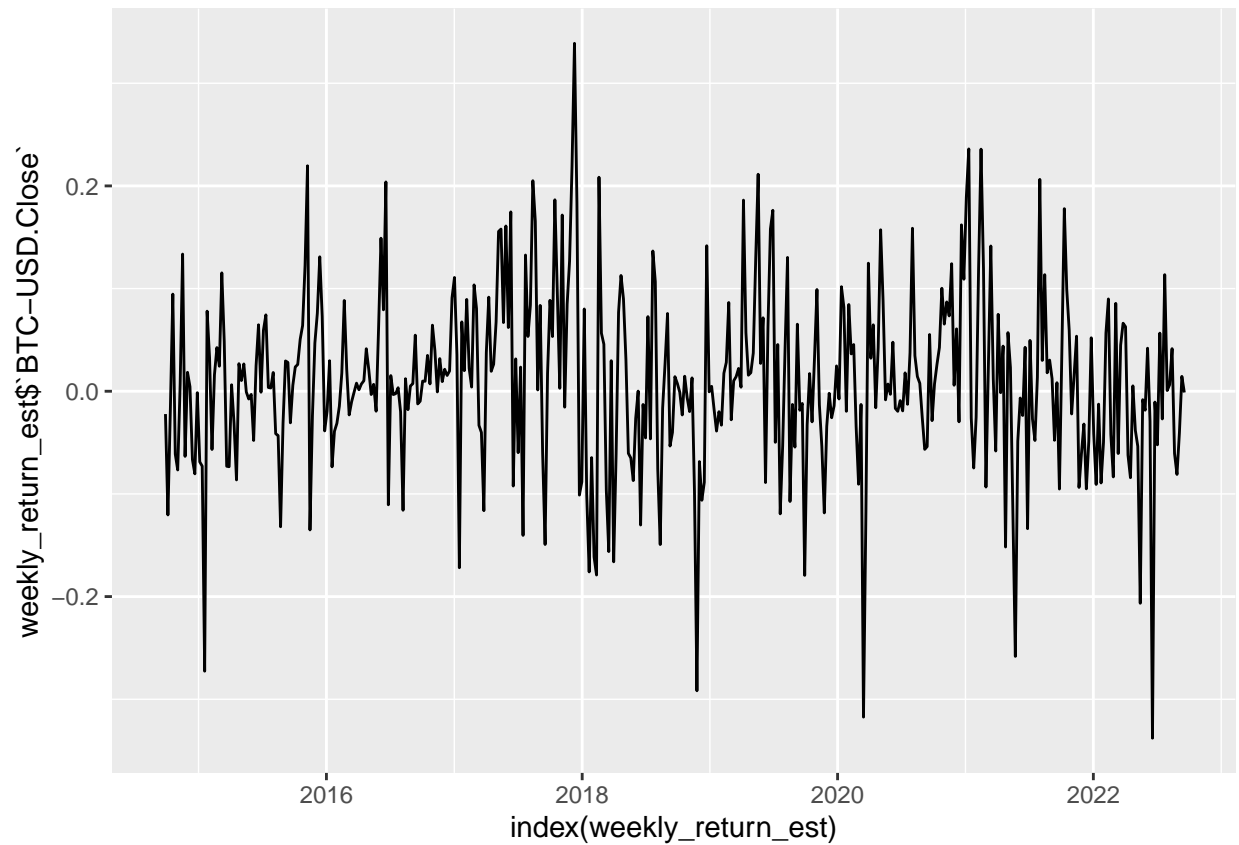


```
weekly_return_est <- weekly_return[index(weekly_return) >= as.Date("2014-09-21") &
  index(weekly_return) <= as.Date("2022-09-21"), ]
weekly_return_hb <- weekly_return[index(weekly_return) >= as.Date("2022-09-22") &
  index(weekly_return) <= as.Date("2024-02-23"), ]
```

```
ggplot(data = weekly_return_est, aes(y = weekly_return_est$`BTC-USD.Close`,
  x = index(weekly_return_est)), group = 1) + geom_line()
```


```
## Don't know how to automatically pick scale for object of type <xts/zoo>.
## Defaulting to continuous.
```

```
## Warning: Removed 1 row containing missing values (`geom_line()`).
```



```
ggplot(data = weekly_return_hb, aes(y = weekly_return_hb$`BTC-USD.Close`,  
  x = index(weekly_return_hb)), group = 1) + geom_line()
```

```
## Don't know how to automatically pick scale for object of type <xts/zoo>.  
## Defaulting to continuous.
```




Model fitting and model selection

Analyzing ACF and PACF

```
acf(weekly_return_est[2:length(weekly_return_est)])
```


Series weekly_return_est[2:length(weekly_return_est)]



We see that the *ACF is smoothly decaying* which means its more *likely to be an AR process*

```
pacf(weekly_return_est[2:length(weekly_return_est)])
```


Series weekly_return_est[2:length(weekly_return_est)]



There is no continuous decay in PACF, the value abruptly falls after first lag. so we can assume it to be a AR process. **Note:** The 19th lag is also significant, suggesting some seasonality. But we will ignore that for now.

Since one lag is significant, AR(1) model can be tested. We will also test ARMA(1,1) and ARMA(1,2) and ARMA(2,1)

Testing ARMA models

AR(1)

```
ar1 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(1,
  0, 0))
summary(ar1)
```

```
##          Length Class  Mode
## coef         2  -none- numeric
## sigma2        1  -none- numeric
## var.coef       4  -none- numeric
## mask          2  -none- logical
## loglik         1  -none- numeric
## aic            1  -none- numeric
## arma           7  -none- numeric
## residuals 418   ts      numeric
```

```
## call      3    -none- call
## series    1    -none- character
## code      1    -none- numeric
## n.cond    1    -none- numeric
## nobs      1    -none- numeric
## model     10   -none- list
```

Coefficients Test The coefficients have some standard error but can be considered insignificant. Statistically testing the significance:

```
library(lmtest)
```

```
## Warning: package 'lmtest' was built under R version 4.2.3
```

```
coeftest(ar1)
```

```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1      0.2807682  0.0469217  5.9838 2.18e-09 ***
## intercept 0.0092963  0.0056354  1.6496 0.09902 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The coefficients are statistically significant. (The Intercept has a lower significance threshold, but can be fitted in for the model.)


```
library("forecast")
```

Residual Check

```
## Warning: package 'forecast' was built under R version 4.2.3
```

```
checkresiduals(ar1)
```

Residuals from ARIMA(1,0,0) with non-zero mean




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,0,0) with non-zero mean
## Q* = 4.7308, df = 9, p-value = 0.8571
##
## Model df: 1.    Total lags used: 10
```

The residuals seem to be normally distributed and have a mean of 0. The ACF has an outlier at 19th lag(as discussed earlier maybe due to seasonality).

Stationarity, Invertibility and Causality Since AR(1), we only need to check for stationarity and causality (only pertaining to $\phi(z)$)

$$\phi(z)y_t = \theta(z)u_t$$

```
autoplot(ar1)
```



roots within the circle, stationary and non-causal.

Mean of the model

$$\mathbb{E}[y_t \sim ARMA(p, q)] = \mathbb{E}[y_t \sim AR(p)] = \frac{a_0}{1 - \sum_{i=1}^p a_i}$$

```
ar1$coef
```

```
##          ar1  intercept
## 0.280768157 0.009296349
```

$$= -0.009/0.2808 \approx 0.3205$$

ARMA(2,1)

```
arma21 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(2,
  0, 1))
arma21
```

```
##
## Call:
```

```
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(2, 0, 1))
##
## Coefficients:

## Warning in sqrt(diag(x$var.coef)): NaNs produced

##          ar1      ar2      ma1  intercept
##          0.1363  0.0495  0.1417      0.0093
## s.e.      NaN      NaN      NaN      0.0057
##
## sigma^2 estimated as 0.006862:  log likelihood = 446.95,  aic = -883.91
```

The coefficients have very high standard error

```
coeftest(arma21)
```

Coefficients Test


```
## Warning in sqrt(diag(se)): NaNs produced

##
## z test of coefficients:
##
##          Estimate Std. Error z value Pr(>|z|)
## ar1          0.1362553      NaN      NaN      NaN
## ar2          0.0495093      NaN      NaN      NaN
## ma1          0.1417153      NaN      NaN      NaN
## intercept 0.0092889  0.0056828  1.6346  0.1021
```

The coefficients are not significant.

```
checkresiduals(arma21)
```


Residuals from ARIMA(2,0,1) with non-zero mean



Residual Check

```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(2,0,1) with non-zero mean
## Q* = 4.5858, df = 7, p-value = 0.7104
##
## Model df: 3.   Total lags used: 10
```

```
autoplot(arma21)
```



Stationarity, Invertibility and Causality

non-causal stationary, non-invertible.

Mean =

$0.0092889 / (1 - (0.1362553 + 0.0495093))$
 = 0.01140812595

ARMA(1,2)

```
arma12 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(1,
  0, 2))
arma12
```

```
##
## Call:
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(1, 0, 2))
##
## Coefficients:
##          ar1          ma1          ma2  intercept
##          0.3042 -0.0259  0.0016      0.0093
## s.e.  1.0178   1.0238  0.2902      0.0057
##
## sigma^2 estimated as 0.006862:  log likelihood = 446.95,  aic = -883.9
```

```
coeftest(arma12)
```


Coefficients Test

```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1      0.3042082  1.0178340  0.2989  0.7650
## ma1     -0.0258956  1.0237949 -0.0253  0.9798
## ma2      0.0015675  0.2901812  0.0054  0.9957
## intercept 0.0092838  0.0056835  1.6335  0.1024
```

Coefficients are not significant.

```
checkresiduals(arma12)
```

Residuals from ARIMA(1,0,2) with non-zero mean




Residual Check

```
##
## Ljung-Box test
##
```



```
## data: Residuals from ARIMA(1,0,2) with non-zero mean
## Q* = 4.6056, df = 7, p-value = 0.708
##
## Model df: 3. Total lags used: 10
```

```
autoplot(arma12)
```



Stationarity, Invertibility and Causality

non-Invertible, stationary and not causal

Mean

```
= 0.0092838/(1-0.3042082)
= 0.01334278443
```

MA(2)

```
ma2 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(0,
  0, 2))
summary(ma2)
```

```
##
## Call:
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(0, 0, 2))
##
## Coefficients:
##          ma1          ma2  intercept
##          0.2777  0.0806      0.0093
## s.e.  0.0487  0.0474      0.0055
##
## sigma^2 estimated as 0.006865:  log likelihood = 446.87,  aic = -885.74
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE          MASE
## Training set 7.18176e-06 0.08285518 0.06097901 105.3459 187.0363 0.7772706
##              ACF1
## Training set 0.001767797
```

Moderately High Standard Errors

```
coeftest(ma2)
```


Coefficients Test

```
##
## z test of coefficients:
##
##          Estimate Std. Error z value Pr(>|z|)
## ma1          0.2776930  0.0486786  5.7046 1.166e-08 ***
## ma2          0.0805563  0.0473753  1.7004  0.08906 .
## intercept 0.0093008  0.0055070  1.6889  0.09124 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The ma1 coefficient appears to be significant that too with a narrower CI, and the ma2 and intercept appear to be significant with loose constraints.

```
checkresiduals(ma2)
```

Residuals from ARIMA(0,0,2) with non-zero mean

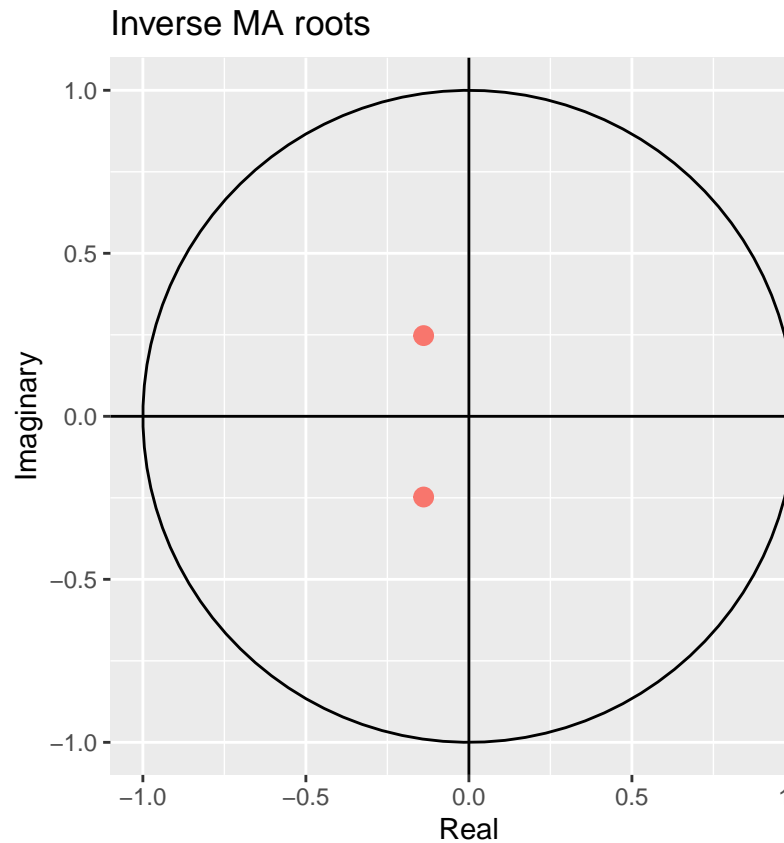


Residual Check

```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,0,2) with non-zero mean
## Q* = 4.7989, df = 8, p-value = 0.7788
##
## Model df: 2.   Total lags used: 10
```

Even the residuals appear to be normally distributed at zero.

```
autoplot(ma2)
```



Stationarity, Causality and Invertibility

There is no invertibility and stationarity is also met. (no causality as well)

Mean = 0

Ljung-Box test for AR(1) MA(2) and ARMA(1,2)

```
print(Box.test(ar1$resid, type = "Ljung-Box", lag = 20))
```

```
##
## Box-Ljung test
##
## data:  ar1$resid
## X-squared = 14.724, df = 20, p-value = 0.792
```

```
print(Box.test(ma2$resid, type = "Ljung-Box", lag = 20))
```


```
##
## Box-Ljung test
##
## data:  ma2$resid
## X-squared = 15.047, df = 20, p-value = 0.7737
```

```
print(Box.test(arma12$resid, type = "Ljung-Box", lag = 20))
```

```
##
## Box-Ljung test
##
## data: arma12$resid
## X-squared = 14.658, df = 20, p-value = 0.7956
```

```
d1 <- checkresiduals(ar1)$statistic
```


Residuals from ARIMA(1,0,0) with non-zero mean



```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(1,0,0) with non-zero mean
## Q* = 4.7308, df = 9, p-value = 0.8571
##
## Model df: 1. Total lags used: 10
```

```
d2 <- checkresiduals(ma2)$statistic
```


Residuals from ARIMA(0,0,2) with non-zero mean



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,0,2) with non-zero mean
## Q* = 4.7989, df = 8, p-value = 0.7788
##
## Model df: 2.   Total lags used: 10
```

```
d3 <- checkresiduals(arma12)$statistic
```

Residuals from ARIMA(1,0,2) with non-zero mean



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,0,2) with non-zero mean
## Q* = 4.6056, df = 7, p-value = 0.708
##
## Model df: 3.   Total lags used: 10
```

```
print(d1)
```

```
##          Q*
## 4.730752
```

```
print(d2)
```

```
##          Q*
## 4.798864
```

```
print(d3)
```

```
##          Q*
## 4.60561
```

```
print(AIC(ar1))
```

AIC

```
## [1] -887.8655
```

```
print(AIC(ma2))
```

```
## [1] -885.7352
```

```
print(AIC(arma12))
```

```
## [1] -883.9013
```

Table 1: Table 1

Col1	ARMA(1,2)	AR(1)	MA(2)
AR Coeff 1	0.3042(1.0178)	0.2808(0.0469)	-
AR Coeff 2	-	-	-
MA Coeff 1	-0.0259(1.0238)	-	0.2777(0.0487)
MA Coeff 2	-0.0016(0.2902)	-	0.0806(0.0474)
AIC	-885.7352	-887.8655	-883.9013
Q-Statistic	4.6056	4.7308	4.7989
p-value of Q-Stat	0.708	0.8571	0.7788

Forecasting

We will forecast for :


```
length(weekly_return_hb)
```

```
## [1] 74
```


74 weeks ahead.

```
ar1_forecast <- forecast(ar1, h = 74)
plot(ar1_forecast)
```


Forecasts from ARIMA(1,0,0) with non-zero mean



```
ts.plot(ar1_forecast$mean, weekly_return_hb, lty = c(1, 2))
```



The forecast seems to not have any significant variation

adding higher order arma models

```
arma192 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(19,
0, 2))
summary(arma192)
```

```
##
## Call:
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(19, 0, 2))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ar6          ar7          ar8          ar9
##      -0.5538  -0.3738   0.1713   0.0248   0.0078   0.0667  -0.0264   0.0312   0.0062
## s.e.    0.2464   0.1509   0.0761   0.0591   0.0596   0.0592   0.0605   0.0610   0.0600
##          ar10         ar11         ar12         ar13         ar14         ar15         ar16         ar17         ar18
##      0.0593   0.0703   0.0505   0.0045  -0.0291  -0.0121   0.0100   0.0386   0.1045
## s.e.   0.0599   0.0593   0.0610   0.0597   0.0604   0.0606   0.0602   0.0608   0.0620
##          ar19         ma1         ma2      intercept
##      -0.0542   0.8476   0.6306           0.009
## s.e.    0.0709   0.2432   0.1762           0.007
##
## sigma^2 estimated as 0.006569:  log likelihood = 455.69,  aic = -865.38
##
## Training set error measures:
```

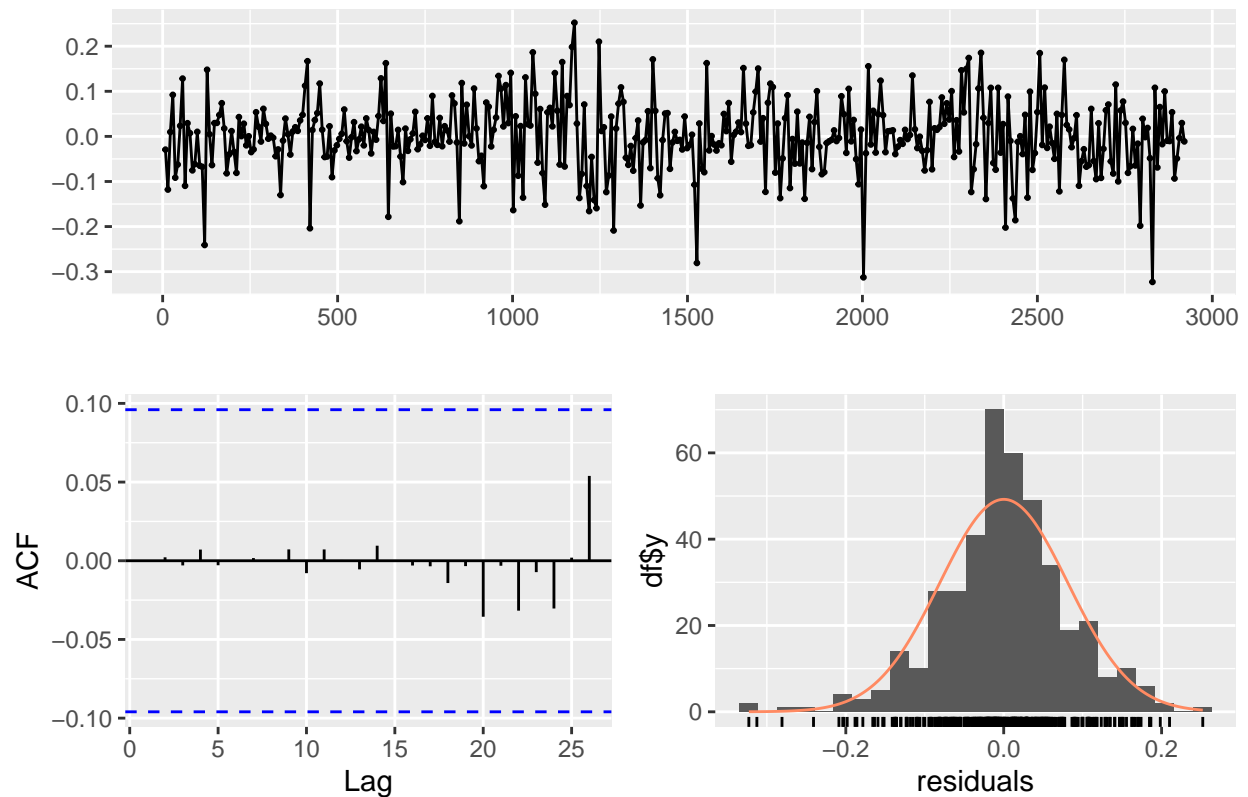
```
##                               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.0001486481 0.08105163 0.06039043 146.0238 241.573 0.7697682
##                               ACF1
## Training set -0.00055882
```

```
coeftest(arma192)
```

```
##
## z test of coefficients:
##
##      Estimate Std. Error z value Pr(>|z|)
## ar1      -0.5537827  0.2463792 -2.2477 0.0245963 *
## ar2      -0.3738090  0.1508605 -2.4778 0.0132178 *
## ar3       0.1713178  0.0760646  2.2523 0.0243054 *
## ar4       0.0248016  0.0591054  0.4196 0.6747654
## ar5       0.0077612  0.0595747  0.1303 0.8963480
## ar6       0.0667446  0.0591619  1.1282 0.2592485
## ar7      -0.0263700  0.0604830 -0.4360 0.6628437
## ar8       0.0311571  0.0609968  0.5108 0.6094919
## ar9       0.0061907  0.0600011  0.1032 0.9178229
## ar10      0.0593488  0.0598620  0.9914 0.3214773
## ar11      0.0703454  0.0593403  1.1855 0.2358368
## ar12      0.0505427  0.0610132  0.8284 0.4074500
## ar13      0.0045300  0.0596643  0.0759 0.9394786
## ar14     -0.0290577  0.0603744 -0.4813 0.6303096
## ar15     -0.0121239  0.0606045 -0.2000 0.8414421
## ar16      0.0099749  0.0602070  0.1657 0.8684114
## ar17      0.0386126  0.0607768  0.6353 0.5252208
## ar18      0.1045358  0.0620363  1.6851 0.0919742 .
## ar19     -0.0541915  0.0709277 -0.7640 0.4448447
## ma1       0.8475674  0.2432463  3.4844 0.0004932 ***
## ma2       0.6306397  0.1762249  3.5786 0.0003454 ***
## intercept 0.0089746  0.0069747  1.2867 0.1981820
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
checkresiduals(arma192)
```

Residuals from ARIMA(19,0,2) with non-zero mean




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(19,0,2) with non-zero mean
## Q* = 1.6973, df = 3, p-value = 0.6375
##
## Model df: 21.    Total lags used: 24
```

Stable residuals with stationarity


Forecasts:

```
arma192_forecast <- forecast(arma192, h = 74)
plot(arma192_forecast)
```

Forecasts from ARIMA(19,0,2) with non-zero mean



```
ts.plot(arma192_forecast$mean, weekly_return_hb$`BTC-USD.Close`,  
        lty = c(1, 2))
```



ARMA(2,19)

```
arma219 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(2,
  0, 19))
summary(arma219)
```

```
##
## Call:
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(2, 0, 19))
##
## Coefficients:
##      ar1      ar2      ma1      ma2      ma3      ma4      ma5      ma6      ma7
##    -1.3653 -0.8123  1.6640  1.3179  0.3835  0.1332  0.0692  0.1082  0.0679
## s.e.   0.1137  0.0863  0.1215  0.1408  0.1190  0.1164  0.1175  0.1183  0.1191
##      ma8      ma9      ma10     ma11     ma12     ma13     ma14     ma15     ma16
##    0.0492  0.0464  0.0937  0.1329  0.1580  0.0935  0.0261  0.0041  0.0134
## s.e.   0.1191  0.1202  0.1248  0.1256  0.1272  0.1273  0.1153  0.1107  0.1172
##      ma17     ma18     ma19 intercept
##    0.0487  0.1484  0.1140      0.009
## s.e.   0.1200  0.1021  0.0566      0.007
##
## sigma^2 estimated as 0.006557:  log likelihood = 455.97,  aic = -865.93
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.000125333 0.08097351 0.06025929 133.7448 233.9468 0.7680967
```


```
##                               ACF1
## Training set -0.0006371101
```

```
coeftest(arma219)
```

```
##
## z test of coefficients:
##
##      Estimate Std. Error  z value  Pr(>|z|)
## ar1      -1.3652974  0.1136869 -12.0093 < 2.2e-16 ***
## ar2      -0.8123236  0.0863371  -9.4087 < 2.2e-16 ***
## ma1       1.6640116  0.1215089  13.6946 < 2.2e-16 ***
## ma2       1.3178631  0.1407999   9.3598 < 2.2e-16 ***
## ma3       0.3834665  0.1190396   3.2213  0.001276 **
## ma4       0.1331676  0.1164381   1.1437  0.252758
## ma5       0.0692331  0.1175255   0.5891  0.555801
## ma6       0.1081623  0.1182639   0.9146  0.360410
## ma7       0.0678771  0.1190634   0.5701  0.568616
## ma8       0.0491956  0.1190705   0.4132  0.679487
## ma9       0.0464292  0.1201987   0.3863  0.699296
## ma10      0.0937248  0.1248442   0.7507  0.452813
## ma11      0.1328901  0.1256324   1.0578  0.290161
## ma12      0.1580267  0.1272333   1.2420  0.214228
## ma13      0.0935150  0.1273053   0.7346  0.462600
## ma14      0.0261148  0.1152997   0.2265  0.820817
## ma15      0.0041462  0.1106546   0.0375  0.970110
## ma16      0.0134102  0.1172433   0.1144  0.908937
## ma17      0.0487361  0.1199793   0.4062  0.684593
## ma18      0.1484105  0.1021240   1.4532  0.146157
## ma19      0.1140386  0.0566172   2.0142  0.043988 *
## intercept 0.0090390  0.0070372   1.2845  0.198981
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
checkresiduals(arma219)
```


Residuals from ARIMA(2,0,19) with non-zero mean




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(2,0,19) with non-zero mean
## Q* = 1.5478, df = 3, p-value = 0.6713
##
## Model df: 21.    Total lags used: 24
```

```
arma219_forecast <- forecast(arma219, h = 74)
plot(arma219_forecast)
```


Forecasts from ARIMA(2,0,19) with non-zero mean



```
ts.plot(arma219_forecast$mean, weekly_return_hb$`BTC-USD.Close`,  
        lty = c(1, 2))
```



```
weekly_return_hb_ts <- as.ts(weekly_return_hb, start = 2927)
errors_ar1 <- weekly_return_hb_ts - ar1_forecast$mean
errors_arma192 <- weekly_return_hb_ts - arma192_forecast$mean
errors_arma219 <- weekly_return_hb_ts - arma219_forecast$mean
```

```
mspear1 <- mean(errors_ar1^2)
mspearma192 <- mean(errors_arma192^2)
mspearma219 <- mean(errors_arma219^2)
```

Paired F-test

```
f_ararma192 <- ((mspear1 - mspearma192)/(74 - 1))/(mspearma192/(74 - 21))
p_value_f <- 1 - pf(f_ararma192, df1 = 74 - 1, df2 = 74 - 21)
p_value_f
```

AR1 ARMA192

```
## [1] 1
```

The p-value of 1 suggests that there is not enough evidence to reject the null hypothesis that the two models have the same Mean Squared Prediction Error (MSPE). A p-value of 1 indicates that the difference in MSPE between the two models is likely due to random chance, rather than a true difference in forecast accuracy.

There might also be a chance of overfit.

```
f_ararma219 <- ((mspear1 - mspearma219)/(74 - 1))/(mspearma219/(74 - 21))
p_value_f <- 1 - pf(f_ararma219, df1 = 74 - 1, df2 = 74 - 21)
p_value_f
```

AR1 ARMA(2,19)

```
## [1] 1
```

Same p-value

```
f_armaarma <- ((mspearma192 - mspearma219)/(74 - 21))/(mspearma219/(74 - 21))
p_value_f <- 1 - pf(f_armaarma, df1 = 74 - 21, df2 = 74 - 21)
p_value_f
```

ARMA(19,2) ARMA(2,19)

```
## [1] 1
```

Same p-value

DM Test

```
dm_test_1 <- dm.test(errors_ar1^2, errors_arma192^2, h = 74)
dm_test_1
```

```
##
## Diebold-Mariano Test
##
## data: errors_ar1^2errors_arma192^2
## DM = 0, Forecast horizon = 74, Loss function power = 2, p-value = 1
## alternative hypothesis: two.sided
```

```
dm_test_2 <- dm.test(errors_ar1^2, errors_arma219^2, h = 74)
```

```
## Warning in dm.test(errors_ar1^2, errors_arma219^2, h = 74): Variance is
## negative. Try varestimator = bartlett. Proceeding with horizon h=1.
```

```
dm_test_2
```

```
##
## Diebold-Mariano Test
##
## data: e1e2
## DM = 0.58644, Forecast horizon = 1, Loss function power = 2, p-value =
## 0.5594
## alternative hypothesis: two.sided
```

```
dm_test_3 <- dm.test(errors_arma192^2, errors_arma219^2, h = 74)
```

```
## Warning in dm.test(errors_arma192^2, errors_arma219^2, h = 74): Variance is
## negative. Try varestimator = bartlett. Proceeding with horizon h=1.
```

```
dm_test_3
```

```
##
## Diebold-Mariano Test
##
## data: e1e2
## DM = 0.58224, Forecast horizon = 1, Loss function power = 2, p-value =
## 0.5622
## alternative hypothesis: two.sided
```

Absolute loss function Since Absolute loss penalizes large errors linearly, whereas quadratic loss penalizes them quadratically, we can have a better differentiation wrt larger errors.

```
dm_test_ar1_arma192 <- dm.test(abs(errors_ar1), abs(errors_arma192),
  alternative = "two.sided", h = 1)
dm_test_ar1_arma219 <- dm.test(abs(errors_ar1), abs(errors_arma219),
  alternative = "two.sided", h = 1)
dm_test_arma192_arma219 <- dm.test(abs(errors_arma192), abs(errors_arma219),
  alternative = "two.sided", h = 1)

# Print DM test results
cat("DM test AR(1) vs ARMA(19,2):", "statistic =", dm_test_ar1_arma192$statistic,
    "p-value =", dm_test_ar1_arma192$p.value, "\n")
```

```
## DM test AR(1) vs ARMA(19,2): statistic = -0.00476356 p-value = 0.9962122
```

```
cat("DM test AR(1) vs ARMA(2,19):", "statistic =", dm_test_ar1_arma219$statistic,
    "p-value =", dm_test_ar1_arma219$p.value, "\n")
```

```
## DM test AR(1) vs ARMA(2,19): statistic = 0.2159198 p-value = 0.8296525
```

```
cat("DM test ARMA(19,2) vs ARMA(2,19):", "statistic =",
    dm_test_arma192_arma219$statistic,
    "p-value =", dm_test_arma192_arma219$p.value, "\n")
```

```
## DM test ARMA(19,2) vs ARMA(2,19): statistic = 0.3085195 p-value = 0.7585658
```

Table 2

	Test Statistic	p-value
AR(1) vs ARMA(19,2)		
F-test	-9.87831e-05	1
DM-test (Quadratic Loss)	0.5157	0.6076
DM-test (Abs Loss)	-0.0047	0.99621
AR(1) vs ARMA(2,19)		
F-test	0.00620222	1
DM-test (Quadratic Loss)	0	1
DM-test (Abs Loss)	0.2154764	0.8299968
AR(19,2) vs ARMA(2,19)		
F-test	0.008679921	1
DM-test (Quadratic Loss)	0.58104	0.563
DM-test (Abs Loss)	0.3078686	0.7590591

```
mean_forecast_ar1 <- mean(ar1_forecast$mean)
mean_forecast_arma192 <- mean(arma192_forecast$mean)
mean_forecast_arma219 <- mean(arma219_forecast$mean)

squared_diff_ar1 <- (ar1_forecast$mean - mean_forecast_ar1)^2
forecast_variance_ar1 <- mean(squared_diff_ar1)

squared_diff_arma192 <- (arma192_forecast$mean - mean_forecast_arma192)^2
forecast_variance_arma192 <- mean(squared_diff_arma192)

squared_diff_arma219 <- (arma219_forecast$mean - mean_forecast_arma219)^2
forecast_variance_arma219 <- mean(squared_diff_arma219)

cat("Forecast Variance AR(1):", forecast_variance_ar1, "\n")
```

```
## Forecast Variance AR(1): 1.246109e-07
```

```
cat("Forecast Variance ARMA(19,2):", forecast_variance_arma192,
    "\n")
```

```
## Forecast Variance ARMA(19,2): 5.636904e-05
```

```
cat("Forecast Variance ARMA(2,19):", forecast_variance_arma219,
    "\n")
```

```
## Forecast Variance ARMA(2,19): 5.907173e-05
```

AR(1) shows least variance but without trend.

Changing the estimation and holdback

```
btc_week_est <- btc_week_df[index(btc_week_df) >= as.Date("2014-09-21") &
  index(btc_week_df) <= as.Date("2023-07-21"), ]

btc_week_hb <- btc_week_df[index(btc_week_df) >= as.Date("2023-07-22") &
  index(btc_week_df) <= as.Date("2024-02-23"), ]
```

80-20

```
weekly_return_est <- weekly_return[index(weekly_return) >= as.Date("2014-09-21") &
  index(weekly_return) <= as.Date("2023-07-21"), ]
weekly_return_hb <- weekly_return[index(weekly_return) >= as.Date("2023-07-22") &
  index(weekly_return) <= as.Date("2024-02-23"), ]
```

```
ar1 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(1,
  0, 0))
print(summary(ar1))
```

```
##
## Call:
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(1, 0, 0))
##
## Coefficients:
##          ar1  intercept
##          0.2757    0.0093
## s.e.    0.0447    0.0052
##
## sigma^2 estimated as 0.006557:  log likelihood = 503.52,  aic = -1001.05
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE          MASE
## Training set 8.84485e-06 0.08097304 0.05889345 150.7741 230.0094 0.7689173
##              ACF1
## Training set 0.001089799
```

```
arma192 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(19,
  0, 2))
print(summary(arma192))
```

```
##
## Call:
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(19, 0, 2))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ar6          ar7          ar8          ar9
```


```
##      -0.5434 -0.4282  0.171  0.0295  0.0078  0.0626 -0.0286  0.0383 -0.0039
## s.e.   0.2228   0.1370  0.070  0.0571  0.0577  0.0571  0.0578  0.0583  0.0575
##      ar10   ar11   ar12   ar13   ar14   ar15   ar16   ar17   ar18
##      0.0572  0.0692  0.0592  0.0098 -0.0135 -0.005 -0.0015  0.0350  0.0838
## s.e.   0.0573  0.0571  0.0585  0.0574  0.0573  0.057  0.0569  0.0577  0.0597
##      ar19   ma1    ma2  intercept
##      -0.0516  0.8344  0.6680    0.0091
## s.e.   0.0638  0.2196  0.1559    0.0063
##
## sigma^2 estimated as 0.006296:  log likelihood = 512.52,  aic = -979.04
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE          MASE
## Training set 0.0001259958 0.07934888 0.05845152 155.078 241.1218 0.7631474
##              ACF1
## Training set -0.0007322027
```

```
arma219 <- arima(weekly_return_est$`BTC-USD.Close`, order = c(2,
0, 19))
print(summary(arma219))
```

```
##
## Call:
## arima(x = weekly_return_est$`BTC-USD.Close`, order = c(2, 0, 19))
##
## Coefficients:
##      ar1      ar2      ma1      ma2      ma3      ma4      ma5      ma6      ma7
##     -1.3595 -0.7962  1.6573  1.2843  0.3517  0.1191  0.0711  0.0981  0.0522
## s.e.   0.1190  0.1039  0.1247  0.1556  0.1148  0.1104  0.1111  0.1114  0.1114
##      ma8      ma9      ma10     ma11     ma12     ma13     ma14     ma15     ma16
##      0.0496  0.0459  0.0801  0.1194  0.1584  0.0968  0.0368  0.0217  0.0149
## s.e.   0.1115  0.1132  0.1183  0.1183  0.1178  0.1180  0.1069  0.1015  0.1102
##      ma17     ma18     ma19  intercept
##      0.0291  0.1199  0.1031    0.0092
## s.e.   0.1139  0.0961  0.0505    0.0064
##
## sigma^2 estimated as 0.006282:  log likelihood = 512.98,  aic = -979.96
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE          MASE
## Training set 0.0001032993 0.07925911 0.05841989 97.92002 263.7854 0.7627345
##              ACF1
## Training set -0.001181196
```


```
arma219_forecast <- forecast(arma219, h = 31)
plot(arma219_forecast)
```

Forecasts from ARIMA(2,0,19) with non-zero mean




```
arma192_forecast <- forecast(arma192, h = 31)  
plot(arma192_forecast)
```


Forecasts from ARIMA(19,0,2) with non-zero mean




```
ar1_forecast <- forecast(ar1, h = 31)  
plot(ar1_forecast)
```


Forecasts from ARIMA(1,0,0) with non-zero mean




```
ts.plot(arma219_forecast$mean, weekly_return_hb, lty = c(1, 2))
```



```
ts.plot(ar1_forecast$mean, weekly_return_hb, lty = c(1, 2))
```



```
ts.plot(arma192_forecast$mean, weekly_return_hb, lty = c(1, 2))
```



```
weekly_return_hb_ts <- as.ts(weekly_return_hb, start = 3228)
errors_ar1 <- weekly_return_hb_ts - ar1_forecast$mean
errors_arma192 <- weekly_return_hb_ts - arma192_forecast$mean
errors_arma219 <- weekly_return_hb_ts - arma219_forecast$mean
mspear1 <- mean(errors_ar1^2)
mspearma192 <- mean(errors_arma192^2)
mspearma219 <- mean(errors_arma219^2)
print(mspear1)
```

```
## [1] 0.002647079
```

```
print(mspearma192)
```

```
## [1] 0.002504827
```

```
print(mspearma219)
```

```
## [1] 0.002514271
```

The error has been reduced but the is still similar (lack of) trend.

```

mean_forecast_ar1 <- mean(ar1_forecast$mean)
mean_forecast_arma192 <- mean(arma192_forecast$mean)
mean_forecast_arma219 <- mean(arma219_forecast$mean)

squared_diff_ar1 <- (ar1_forecast$mean - mean_forecast_ar1)^2
forecast_variance_ar1 <- mean(squared_diff_ar1)

squared_diff_arma192 <- (arma192_forecast$mean - mean_forecast_arma192)^2
forecast_variance_arma192 <- mean(squared_diff_arma192)

squared_diff_arma219 <- (arma219_forecast$mean - mean_forecast_arma219)^2
forecast_variance_arma219 <- mean(squared_diff_arma219)

cat("Forecast Variance AR(1):", forecast_variance_ar1, "\n")

```

```
## Forecast Variance AR(1): 1.032947e-07
```

```

cat("Forecast Variance ARMA(19,2):", forecast_variance_arma192,
    "\n")

```

```
## Forecast Variance ARMA(19,2): 6.290917e-05
```

```

cat("Forecast Variance ARMA(2,19):", forecast_variance_arma219,
    "\n")

```

```
## Forecast Variance ARMA(2,19): 3.653699e-05
```

AR1 returns minimum forecast variance, but doesn't show any trend.

Despite inaccuracy due to seasonality of data we do have consistent results.